

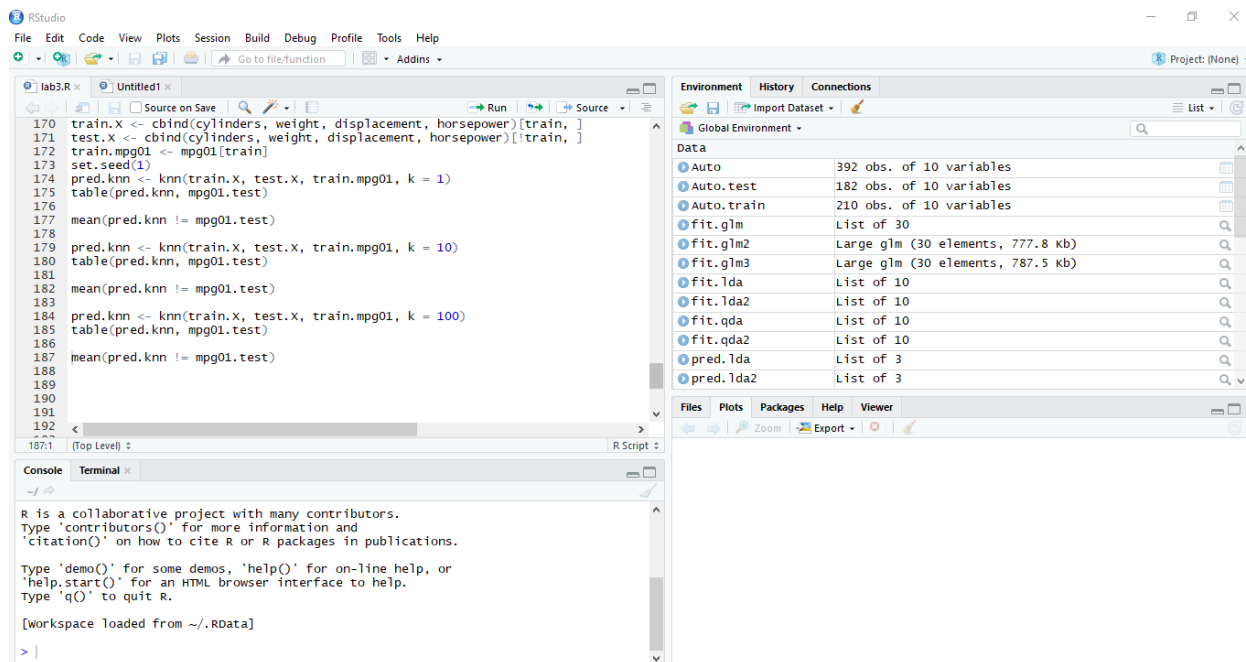
Introduction to Statistical Learning Lab3

Name: Sandeep Reddy Salkuti

Id: 16296868

Email: sswf7@umsystem.edu

1. You may download the R Code for Labs and the Data Sets to use from the textbook website.



The screenshot shows the RStudio environment. The source editor on the left contains R code for a k-NN classification task. The Global Environment pane on the right lists the objects created in the workspace. The console at the bottom shows the R startup message.

```
170 train.X <- cbind(cylinders, weight, displacement, horsepower)[train, ]
171 test.X <- cbind(cylinders, weight, displacement, horsepower)[!train, ]
172 train.mpg01 <- mpg01[train, ]
173 set.seed(1)
174 pred.knn <- knn(train.X, test.X, train.mpg01, k = 1)
175 table(pred.knn, mpg01.test)
176
177 mean(pred.knn != mpg01.test)
178
179 pred.knn <- knn(train.X, test.X, train.mpg01, k = 10)
180 table(pred.knn, mpg01.test)
181
182 mean(pred.knn != mpg01.test)
183
184 pred.knn <- knn(train.X, test.X, train.mpg01, k = 100)
185 table(pred.knn, mpg01.test)
186
187 mean(pred.knn != mpg01.test)
188
189
190
191
192
```

Global Environment

Object	Details
Auto	392 obs. of 10 variables
Auto.test	182 obs. of 10 variables
Auto.train	210 obs. of 10 variables
fit.glm	List of 30
fit.glm2	Large glm (30 elements, 777.8 kb)
fit.glm3	Large glm (30 elements, 787.5 kb)
fit.lda	List of 10
fit.lda2	List of 10
fit.qda	List of 10
fit.qda2	List of 10
pred.lda	List of 3
pred.lda2	List of 3

Console

```
R is a collaborative project with many contributors.
Type 'contributors()' for more information and
'citation()' on how to cite R or R packages in publications.

Type 'demo()' for some demos, 'help()' for on-line help, or
'help.start()' for an HTML browser interface to help.
Type 'q()' to quit R.

[workspace loaded from ~/.RData]
>
```

2. This question should be answered using the Weekly data set, which is part of the ISLR package. This data is similar in nature to the Smarket data from this chapters lab, except that it contains 1,089 weekly returns for 21 years, from the beginning of 1990 to the end of 2010.

(a) (5 points) Produce some numerical and graphical summaries of the Weekly data. Do there appear to be any patterns?

RStudio

File Edit Code View Plots Session Build Debug Profile Tools Help

Go to file/function Addins

Untitled1*

Source on Save Run Source

```
1 #3. a) numerical summaries of weekly data
2 library(ISLR)
3 summary(weekly)
4
```

1:1 (Top Level) R Script

Console Terminal

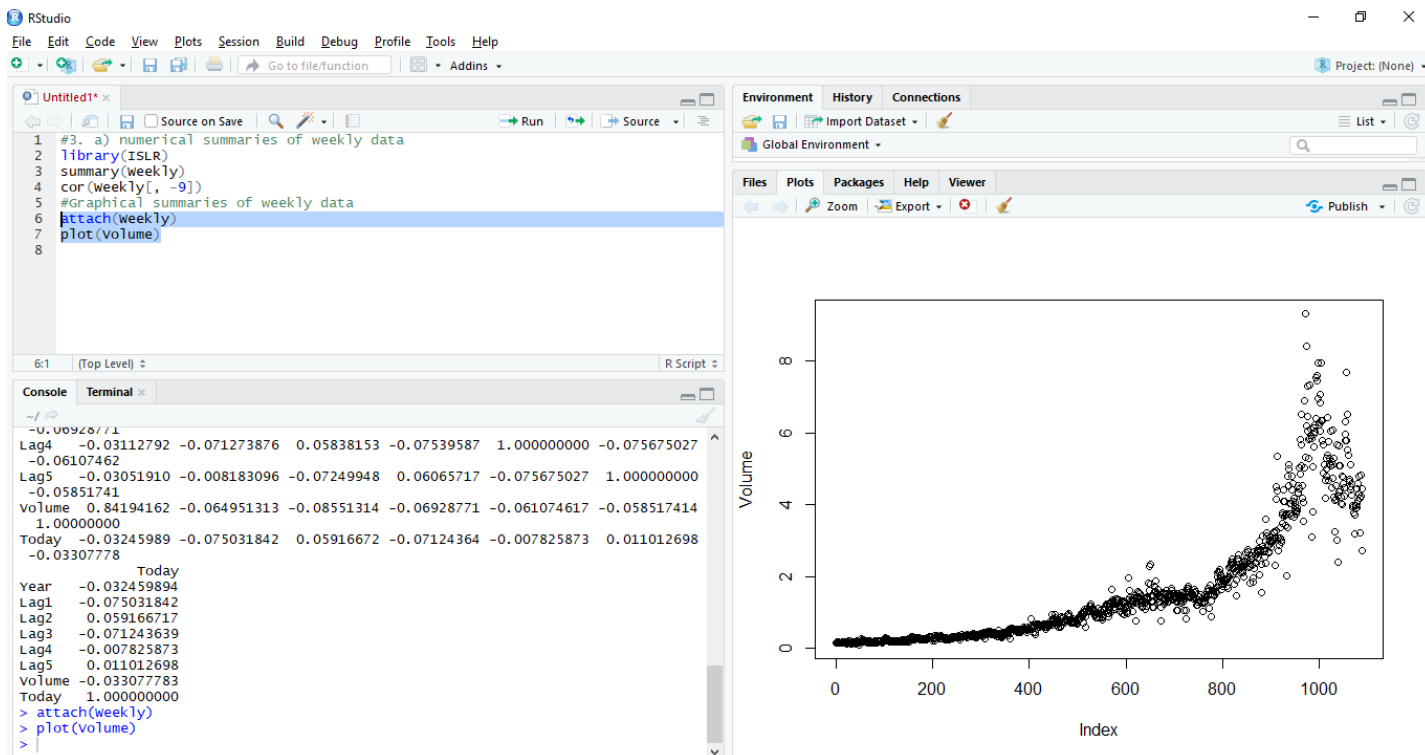
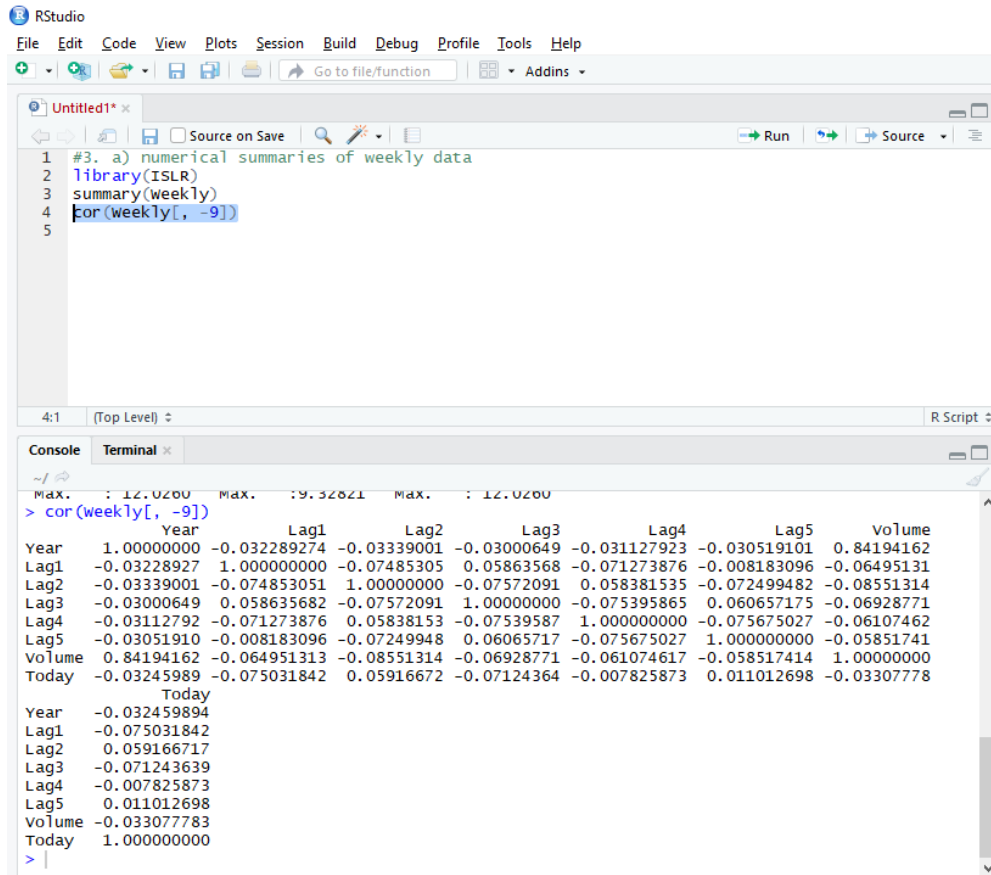
~/

'help.start()' for an HTML browser interface to help.
Type 'q()' to quit R.

```
> library(ISLR)
warning message:
package 'ISLR' was built under R version 3.6.3
> summary(weekly)
```

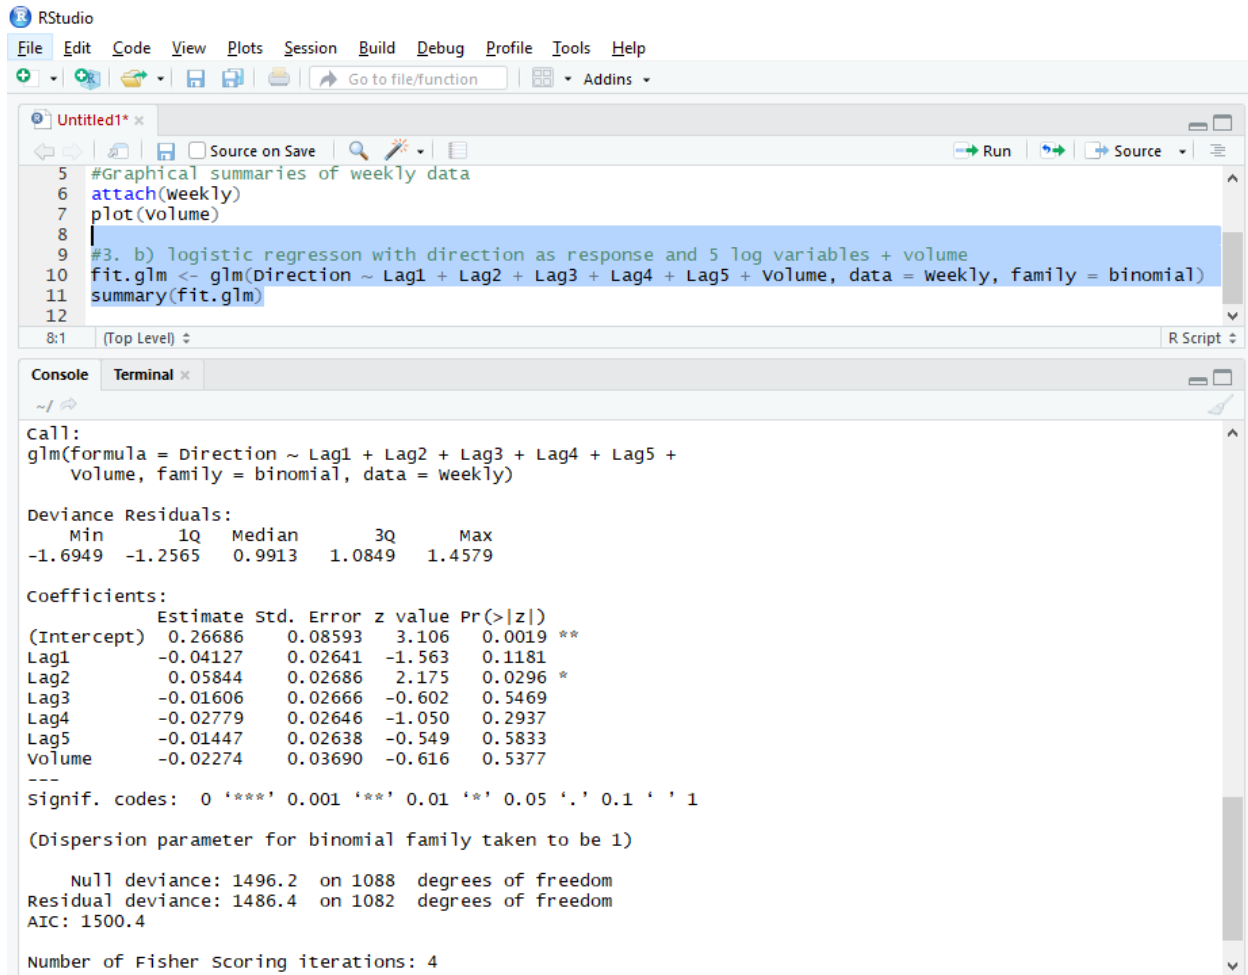
Year		Lag1		Lag2		Lag3		Lag4	
Min.	:1990	Min.	:-18.1950	Min.	:-18.1950	Min.	:-18.1950	Min.	:-18.1950
1st Qu.	:1995	1st Qu.	:-1.1540	1st Qu.	:-1.1540	1st Qu.	:-1.1580	1st Qu.	:-1.1580
Median	:2000	Median	: 0.2410	Median	: 0.2410	Median	: 0.2410	Median	: 0.2380
Mean	:2000	Mean	: 0.1506	Mean	: 0.1511	Mean	: 0.1472	Mean	: 0.1458
3rd Qu.	:2005	3rd Qu.	: 1.4050	3rd Qu.	: 1.4090	3rd Qu.	: 1.4090	3rd Qu.	: 1.4090
Max.	:2010	Max.	: 12.0260	Max.	: 12.0260	Max.	: 12.0260	Max.	: 12.0260

Lag5		Volume		Today		Direction	
Min.	:-18.1950	Min.	:0.08747	Min.	:-18.1950	Down	:484
1st Qu.	:-1.1660	1st Qu.	:0.33202	1st Qu.	:-1.1540	Up	:605
Median	: 0.2340	Median	:1.00268	Median	: 0.2410		
Mean	: 0.1399	Mean	:1.57462	Mean	: 0.1499		
3rd Qu.	: 1.4050	3rd Qu.	:2.05373	3rd Qu.	: 1.4050		
Max.	: 12.0260	Max.	:9.32821	Max.	: 12.0260		



From numerical summaries correlations between lag variables and today's returns are close to zero. There exists substantial correlation between "year" and "volume". From the above plot it is clear that volume increases for index values and

(b) (5 points) Use the full data set to perform a logistic regression with Direction as the response and the five lag variables plus Volume as predictors. Use the summary function to print the results. Do any of the predictors appear to be statistically significant? If so, which ones?



The screenshot shows the RStudio interface. The script editor contains the following code:

```

5 #Graphical summaries of weekly data
6 attach(weekly)
7 plot(volume)
8
9 #3. b) logistic regression with direction as response and 5 lag variables + volume
10 fit.glm <- glm(Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 + volume, data = weekly, family = binomial)
11 summary(fit.glm)
12

```

The console output shows the results of the logistic regression:

```

Call:
glm(formula = Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 +
    volume, family = binomial, data = weekly)

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-1.6949  -1.2565   0.9913   1.0849   1.4579

Coefficients:
            Estimate Std. Error z value Pr(>|z|)
(Intercept)  0.26686    0.08593   3.106  0.0019 **
Lag1        -0.04127    0.02641  -1.563  0.1181
Lag2         0.05844    0.02686   2.175  0.0296 *
Lag3        -0.01606    0.02666  -0.602  0.5469
Lag4        -0.02779    0.02646  -1.050  0.2937
Lag5        -0.01447    0.02638  -0.549  0.5833
volume      -0.02274    0.03690  -0.616  0.5377
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

    Null deviance: 1496.2  on 1088  degrees of freedom
Residual deviance: 1486.4  on 1082  degrees of freedom
AIC: 1500.4

Number of Fisher Scoring iterations: 4

```

From the above it is clear that "Lag2" appears to be only predictor that is statistically significant. It is because of its p value is less than 0.05 ($\Pr(>|z|) = 3\%$.)

(c) (5 points) Compute the confusion matrix and overall fraction of correct predictions. Explain what the confusion matrix is telling you about the types of mistakes made by logistic regression.

```
12 |
13 #3. c) confusion matrix and overall fraction of correct predictions
14 probs <- predict(fit.glm, type = "response")
15 pred.glm <- rep("Down", length(probs))
16 pred.glm[probs > 0.5] <- "up"
17 table(pred.glm, Direction)
```

12:1 (Top Level) R Script

Console Terminal x

```
> probs <- predict(fit.glm, type = "response")
> pred.glm <- rep("Down", length(probs))
> pred.glm[probs > 0.5] <- "up"
> table(pred.glm, Direction)
      Direction
pred.glm Down Up
Down     54  48
Up      430 557
> |
```

From above it is clear that percentage of correct predictions on the training data is $(54 + 557)/1089 = 0.5610$ which is equal to 56.1%. In contrast training error rate is 43.89% which is often overly optimistic. We can say that for weeks when the market goes up, the model is 92.0661% of the time $(557/(48+557))$. For weeks when the market goes down, the model is only 11.15702% of the time $(54/(54+430))$.

(d) (5 points) Now fit the logistic regression model using a training data period from 1990 to 2008, with Lag2 as the only predictor. Compute the confusion matrix and the overall fraction of correct predictions for the held out data (that is, the data from 2009 and 2010).

```

RStudio
File Edit Code View Plots Session Build Debug Profile Tools Help
Go to file/function Addins
Untitled1*
18
19 #3. d) Now fit the logistic regression model using a training data period from 1990 to 2008, with "Lag2"
20 #as the only predictor. Compute the confusion matrix and the overall fraction of correct predictions for
21 #the held out data (that is, the data from 2009 to 2010).
22 train <- (Year < 2009)
23 weekly.20092010 <- weekly[!train, ]
24 Direction.20092010 <- Direction[!train]
25 fit.glm2 <- glm(Direction ~ Lag2, data = weekly, family = binomial, subset = train)
26 summary(fit.glm2)
22:1 (Top Level) R Script
Console Terminal
~/
call:
glm(formula = Direction ~ Lag2, family = binomial, data = weekly,
    subset = train)

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-1.536  -1.264   1.021   1.091   1.368

Coefficients:
              Estimate Std. Error z value Pr(>|z|)
(Intercept)  0.20326    0.06428   3.162  0.00157 **
Lag2         0.05810    0.02870   2.024  0.04298 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

    Null deviance: 1354.7  on 984  degrees of freedom
Residual deviance: 1350.5  on 983  degrees of freedom
AIC: 1354.5

Number of Fisher Scoring iterations: 4
> |

```

```

27
28 probs2 <- predict(fit.glm2, weekly.20092010, type = "response")
29 pred.glm2 <- rep("Down", length(probs2))
30 pred.glm2[probs2 > 0.5] <- "Up"
31 table(pred.glm2, Direction.20092010)
32
27:1 (Top Level)
Console Terminal
~/

```

```

> probs2 <- predict(fit.glm2, weekly.20092010, type = "response")
> pred.glm2 <- rep("Down", length(probs2))
> pred.glm2[probs2 > 0.5] <- "Up"
> table(pred.glm2, Direction.20092010)
      Direction.20092010
pred.glm2 Down Up
Down      9  5
Up      34 56
> |

```

From the above it is clear that the percentage of correct predictions on the test data is $(9+56)/104$ which equals to 62.5%. In other words test error rate is 37.5%. Also for weeks when the market goes up, the

model is 91.803% of the time ($56/(56+5)$). For weeks when the market goes down, the model is only 20.930% of the time ($9/(9+34)$).

(e) (5 points) Repeat (d) using LDA.

```

32
33 #3. e) Repeat (d) using LDA.
34 library(MASS)
35 fit.lda <- lda(Direction ~ Lag2, data = weekly, subset = train)
36 fit.lda
37
38
39
40 <
33:1 (Top Level)

```

```

Console Terminal x
~/
> library(MASS)
Warning message:
package 'MASS' was built under R version 3.6.3
> fit.lda <- lda(Direction ~ Lag2, data = weekly, subset = train)
> fit.lda
Call:
lda(Direction ~ Lag2, data = weekly, subset = train)

Prior probabilities of groups:
      Down      Up
0.4477157 0.5522843

Group means:
      Lag2
Down -0.03568254
Up    0.26036581

Coefficients of linear discriminants:
      LD1
Lag2 0.4414162

```

```

57
58
59 pred.lda <- predict(fit.lda, weekly.20092010)
60 table(pred.lda$class, Direction.20092010)
61
62 <
58:1 (Top Level)

```

```

Console Terminal x
~/
> pred.lda <- predict(fit.lda, weekly.20092010)
> table(pred.lda$class, Direction.20092010)
      Direction.20092010
      Down Up
Down      9  5
Up      34 56
>

```

From above it is clear that percentage of correct predictions on the test data is 62.5%. In other words test error rate is 37.5%. We can say that for weeks when the market goes up, the model is 91.803% of the time. For weeks when the market goes down, the model is only 20.9302% of the time. These results are very close to those obtained with the logistic regression model which is not surprising.

(f) (5 points) Repeat (d) using QDA.

```
40 #3. f) Repeat (d) using QDA.
41 fit.qda <- qda(Direction ~ Lag2, data = weekly, subset = train)
42 fit.qda
43
44
45
46
47
48
49 <
```

40:1 (Top Level) ↕

Console Terminal x

~/

```
> fit.qda <- qda(Direction ~ Lag2, data = weekly, subset = train)
> fit.qda
Call:
qda(Direction ~ Lag2, data = weekly, subset = train)

Prior probabilities of groups:
      Down      Up 
0.4477157 0.5522843 

Group means:
      Lag2
Down -0.03568254
Up    0.26036581
> |
```

```
57 #3. f) Repeat (d) using QDA.
58 fit.qda <- qda(Direction ~ Lag2, data = weekly, subset = train)
59 fit.qda
60
61
62 pred.qda <- predict(fit.qda, weekly.20092010)
63 table(pred.qda$class, Direction.20092010)
64
65 <
```

62:1 (Top Level) ↕

Console Terminal x

~/

```
> pred.qda <- predict(fit.qda, weekly.20092010)
> table(pred.qda$class, Direction.20092010)
      Direction.20092010
      Down Up
Down      0  0
Up       43 61
> |
```

From above it is clear that the percentage of correct predictions on the test data is 58.6538%. In other words test error rate is 41.3461%. We could also say that for weeks when the market goes up, the model is 100% of the time. For weeks when the market goes down, the model is only 0% of the time. It

is clear that QDA achieves a correctness of 58.6538% even though the model chooses “Up” the whole time.

(g) (5 points) Repeat (d) using KNN with K = 1.

```
47 #3. g) Repeat (d) using KNN with K = 1.
48 library(class)
49 train.X <- as.matrix(Lag2[train])
50 test.X <- as.matrix(Lag2[!train])
51 train.Direction <- Direction[train]
52 set.seed(1)
53 pred.knn <- knn(train.X, test.X, train.Direction, k = 1)
54 table(pred.knn, Direction.20092010)
55
56
57 <
```

48:1 (Top Level) ▾

Console Terminal x

~/

C:\Users\SandeepReddy\AppData\Local\Temp\RtmpwnA9no\downloaded_packages

```
> library(class)
warning message:
package 'class' was built under R version 3.6.3
> train.X <- as.matrix(Lag2[train])
> test.X <- as.matrix(Lag2[!train])
> train.Direction <- Direction[train]
> set.seed(1)
> pred.knn <- knn(train.X, test.X, train.Direction, k = 1)
> table(pred.knn, Direction.20092010)
      Direction.20092010
pred.knn Down Up
Down      21  30
Up        22  31
> |
```

From above it is clear that the percentage of correct predictions on the test data is 50%. In other words test error rate is 50%. We can also say that for weeks when the market goes up, the model is 50.8196% of the time. For weeks when the market goes down, the model is only 48.8372% of the time.

(h) (5 points) Which of these methods appears to provide the best results on this data?

After comparison of all test error rates, It is clear that logistic regression and LDA have the minimum error rates, followed by QDA and KNN.

- (i) (5 points) Experiment with different combinations of predictors, including possible transformations and interactions, for each of the methods. Report the variables, method, and associated confusion matrix that appears to provide the best results on the held out data. Note that you should also experiment with values for K in the KNN classifier.**

Logistic regression with Lag2:Lag1

```
56 #3. i)Experiment with different combinations of predictors,  
57 #including possible transformations and interactions, for each of the  
58 #methods. Report the variables, method, and associated confusion matrix  
59 #that appears to provide the best results on the held out data. Note that  
60 #you should also experiment with values for K in the KNN classifier.  
61  
62 #logistic regression with Lag2:Lag1  
63 fit.glm3 <- glm(Direction ~ Lag2:Lag1, data = weekly, family = binomial, subset = train)  
64 probs3 <- predict(fit.glm3, weekly.20092010, type = "response")  
65 pred.glm3 <- rep("Down", length(probs3))  
66 pred.glm3[probs3 > 0.5] = "Up"  
67 table(pred.glm3, Direction.20092010)
```

```
68  
69  
70  
71  
72  
73  
74  
75  
76  
77  
78  
79 < [Top Level] ⚡  
63:1 (Top Level) ⚡  
Console Terminal x  
~/  
> fit.glm3 <- glm(Direction ~ Lag2:Lag1, data = weekly, family = binomial, subset = train)  
> probs3 <- predict(fit.glm3, weekly.20092010, type = "response")  
> pred.glm3 <- rep("Down", length(probs3))  
> pred.glm3[probs3 > 0.5] = "Up"  
> table(pred.glm3, Direction.20092010)  
      Direction.20092010  
pred.glm3 Down Up  
Down      1  1  
Up       42 60  
> |  
  
> mean(pred.glm3 == Direction.20092010)  
[1] 0.5865385  
> |
```

LDA with Lag2 interaction with Lag1

```
> # LDA with Lag2 interaction with Lag1  
> fit.llda2 <- lda(Direction ~ Lag2:Lag1, data = weekly, subset = train)  
> pred.llda2 <- predict(fit.llda2, weekly.20092010)  
> mean(pred.llda2$class == Direction.20092010)  
[1] 0.5769231  
> |
```

QDA with $\sqrt{\text{abs}(\text{Lag2})}$

```
> # QDA with sqrt(abs(Lag2))
> fit.qda2 <- qda(Direction ~ Lag2 + sqrt(abs(Lag2)), data = weekly, subset = train)
> pred.qda2 <- predict(fit.qda2, weekly.20092010)
> table(pred.qda2$class, Direction.20092010)
      Direction.20092010
      Down Up
Down      12 13
Up       31 48
> |
```

```
> mean(pred.qda2$class == Direction.20092010)
[1] 0.5769231
> |
```

KNN with K=10

```
> # KNN k =10
> pred.knn2 <- knn(train.X, test.X, train.Direction, k = 10)
> table(pred.knn2, Direction.20092010)
      Direction.20092010
pred.knn2 Down Up
      Down   17 18
      Up    26 43
> |
```

```
> mean(pred.knn2 == Direction.20092010)
[1] 0.5769231
> |
```

KNN with K=100

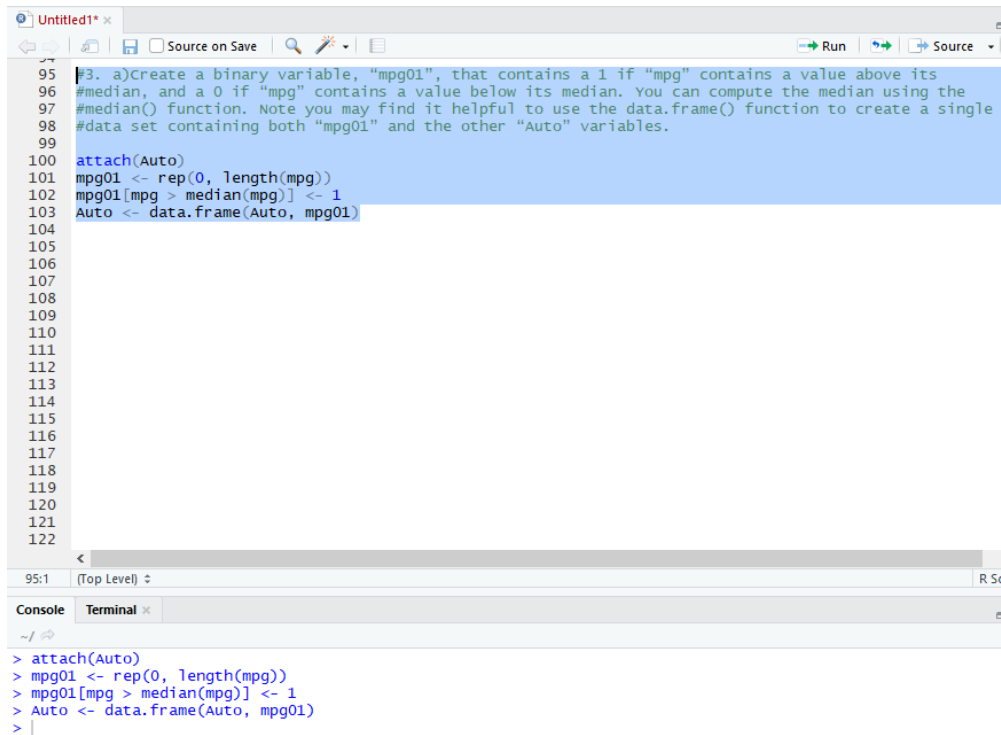
```
> # KNN k = 100
> pred.knn3 <- knn(train.X, test.X, train.Direction, k = 100)
> table(pred.knn3, Direction.20092010)
      Direction.20092010
pred.knn3 Down Up
      Down    9 12
      Up    34 49
> |
```

```
> mean(pred.knn3 == Direction.20092010)
[1] 0.5576923
> |
```

Out of all of these combinations, the original logistic regression and LDA have the best performance in terms of test error rates.

3. In this problem, you will develop a model to predict whether a given car gets high or low gas mileage based on the Auto data set.

(a) (5 points) Create a binary variable, `mpg01`, that contains a 1 if `mpg` contains a value above its median, and a 0 if `mpg` contains a value below its median. You can compute the median using the `median()` function. Note you may find it helpful to use the `data.frame()` function to create a single data set containing both `mpg01` and the other Auto variables.



```
95 #3. a) create a binary variable, "mpg01", that contains a 1 if "mpg" contains a value above its
96 #median, and a 0 if "mpg" contains a value below its median. You can compute the median using the
97 #median() function. Note you may find it helpful to use the data.frame() function to create a single
98 #data set containing both "mpg01" and the other "Auto" variables.
99
100 attach(Auto)
101 mpg01 <- rep(0, length(mpg))
102 mpg01[mpg > median(mpg)] <- 1
103 Auto <- data.frame(Auto, mpg01)
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
```

95:1 (Top Level) R Sc

Console Terminal x

```
> attach(Auto)
> mpg01 <- rep(0, length(mpg))
> mpg01[mpg > median(mpg)] <- 1
> Auto <- data.frame(Auto, mpg01)
>
```

(b) (5 points) Explore the data graphically in order to investigate the association between mpg01 and the other features. Which of the other features seem most likely to be useful in predicting mpg01? Scatterplots and boxplots may be useful tools to answer this question. Describe your findings.

```

104
105 #3. b) Explore the data graphically in order to investigate the association between "mpg01" and the
106 #other features. Which of the other features seem most likely to be useful in predicting "mpg01"?
107 #Scatterplots and boxplots may be useful tools to answer this question. Describe your findings.
108
109 cor(Auto[, -9])
110
111
112
113
114
115
116
105:1 (Top Level) R

```

```

Console Terminal
~/
> cor(Auto[, -9])
      mpg cylinders displacement horsepower weight acceleration year origin
mpg    1.0000000 -0.7776175  -0.8051269  -0.7784268  -0.8322442  0.4233285  0.5805410  0.5652088
cylinders -0.7776175  1.0000000   0.9508233   0.8429834   0.8975273  -0.5046834  -0.3456474  -0.5689316
displacement -0.8051269  0.9508233   1.0000000   0.8972570   0.9329944  -0.5438005  -0.3698552  -0.6145351
horsepower -0.7784268  0.8429834   0.8972570   1.0000000   0.8645377  -0.6891955  -0.4163615  -0.4551715
weight -0.8322442  0.8975273   0.9329944   0.8645377   1.0000000  -0.4168392  -0.3091199  -0.5850054
acceleration 0.4233285 -0.5046834  -0.5438005  -0.6891955  -0.4168392  1.0000000  0.2903161  0.2127458
year 0.5805410 -0.3456474  -0.3698552  -0.4163615  -0.3091199  0.2903161  1.0000000  0.1815277
origin 0.5652088 -0.5689316  -0.6145351  -0.4551715  -0.5850054  0.2127458  0.1815277  1.0000000
mpg01 0.8369392 -0.7591939  -0.7534766  -0.6670526  -0.7577566  0.3468215  0.4299042  0.5136984

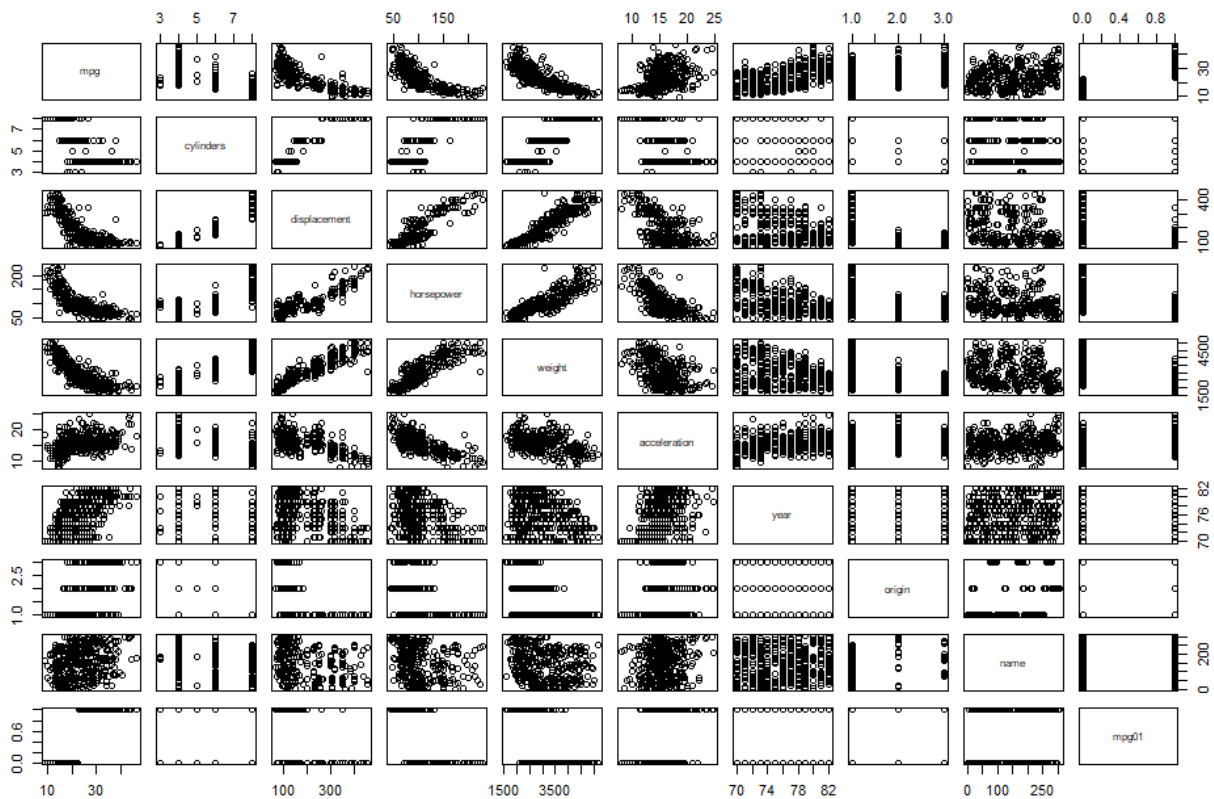
      mpg01
mpg    0.8369392
cylinders -0.7591939
displacement -0.7534766
horsepower -0.6670526
weight -0.7577566
acceleration 0.3468215
year 0.4299042
origin 0.5136984
mpg01 1.0000000
>

```

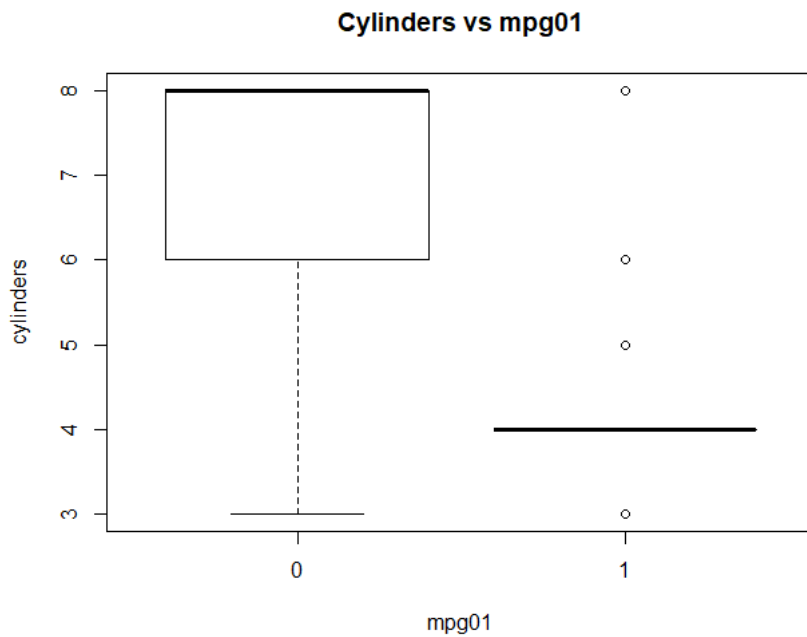
```

> pairs(Auto)
>


```




```
> boxplot(cylinders ~ mpg01, data = Auto, main = "Cylinders vs mpg01")
> |
```



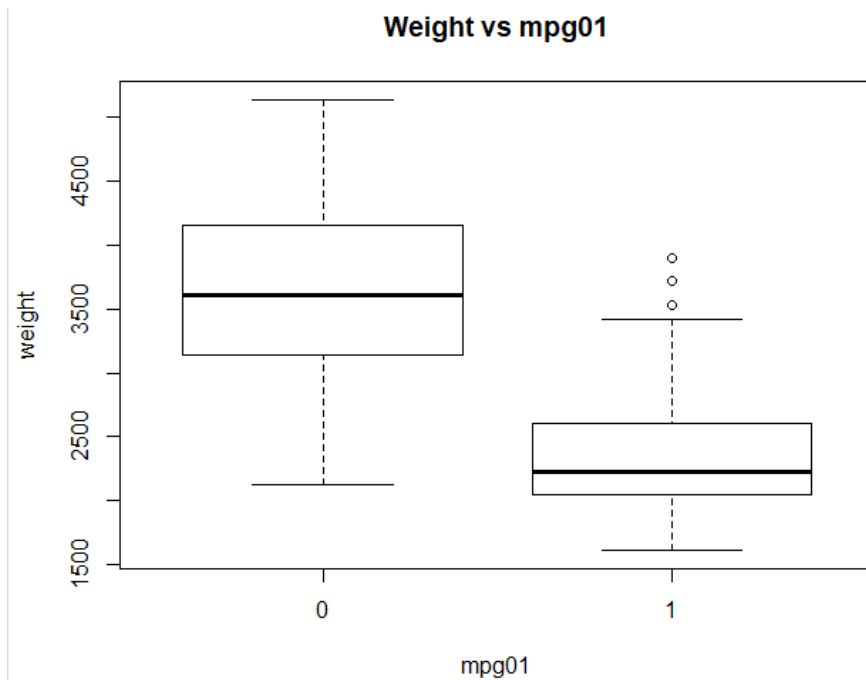
```
> boxplot(displacement ~ mpg01, data = Auto, main = "Displacement vs mpg01")
> |
```



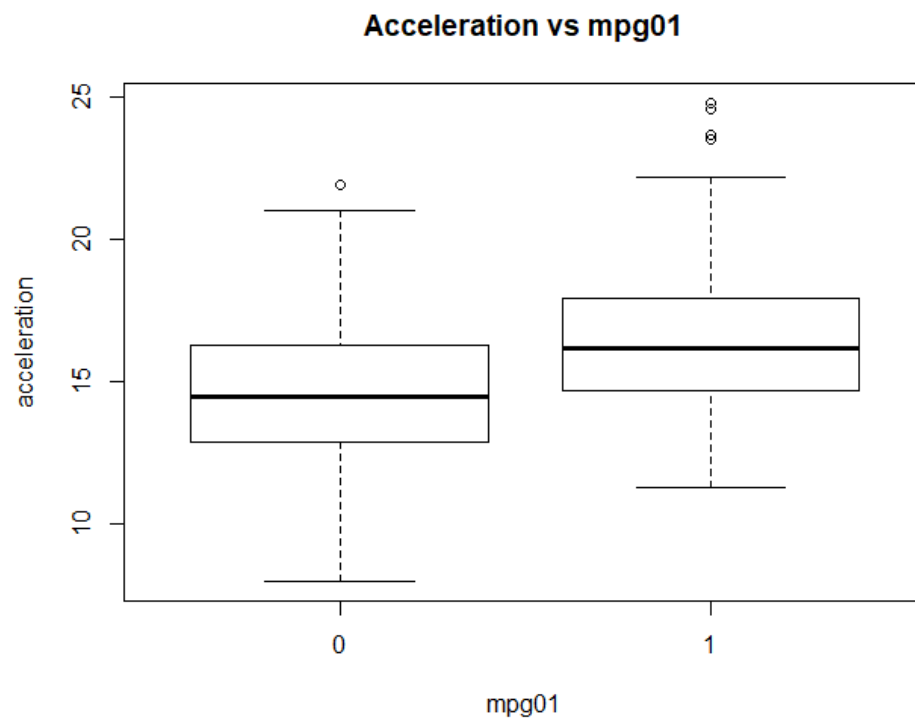
```
> boxplot(horsepower ~ mpg01, data = Auto, main = "Horsepower vs mpg01")
> |
```



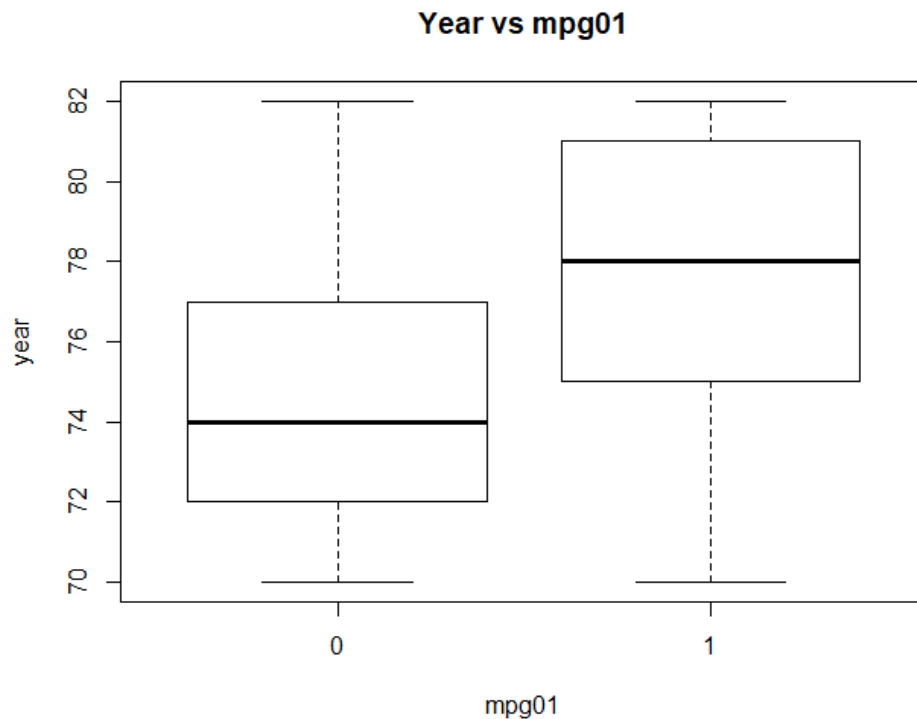
```
> boxplot(weight ~ mpg01, data = Auto, main = "weight vs mpg01")
> |
```



```
> boxplot(acceleration ~ mpg01, data = Auto, main = "Acceleration vs mpg01")  
> |
```



```
> boxplot(year ~ mpg01, data = Auto, main = "Year vs mpg01")  
> |
```

For predicting mpg01 I used box plot for all data against mpg01 and found that there exists some association between “mpg01” and “cylinders”, “weight”, “displacement” and “horsepower”.

(c) Split the data into a training set and a test set.

```
> #3. c) Split the data into a training set and a test set.  
> train <- (year %% 2 == 0)  
> Auto.train <- Auto[train, ]  
> Auto.test <- Auto[!train, ]  
> mpg01.test <- mpg01[!train]  
> |
```

(d) (5 points) Perform LDA on the training data in order to predict mpg01 using the variables that seemed most associated with mpg01 in (b). What is the test error of the model obtained?

```
130
131 #3. d) Perform LDA on the training data in order to predict mpg01 using the
132 #variables that seemed most associated with mpg01 in (b). What is the test error
133 #of the model obtained?
134
135 fit.lda <- lda(mpg01 ~ cylinders + weight + displacement + horsepower, data = Auto, subset = train)
136 fit.lda
137
138
139
140 <
131:1 (Top Level)
R Script

Console Terminal
~/
> fit.lda <- lda(mpg01 ~ cylinders + weight + displacement + horsepower, data = Auto, subset = train)
> fit.lda
call:
lda(mpg01 ~ cylinders + weight + displacement + horsepower, data = Auto,
    subset = train)

Prior probabilities of groups:
      0      1
0.4571429 0.5428571

Group means:
  cylinders    weight displacement horsepower
0  6.812500 3604.823    271.7396   133.14583
1  4.070175 2314.763    111.6623    77.92105

Coefficients of linear discriminants:
              LD1
cylinders    -0.6741402638
weight       -0.0011465750
displacement  0.0004481325
horsepower    0.0059035377
>

> pred.lda <- predict(fit.lda, Auto.test)
> table(pred.lda$class, mpg01.test)
      mpg01.test
      0      1
0  86      9
1  14     73

>
> mean(pred.lda$class != mpg01.test)
[1] 0.1263736
>
```

From above it is clear that the test error rate is 12.637%

(e) (5 points) Perform QDA on the training data in order to predict mpg01 using the variables that seemed most associated with mpg01 in (b). What is the test error of the model obtained?

```
> #3. e) Perform QDA on the training data in order to predict "mpg01" using the variables that seemed
> #most associated with "mpg01" in (b). What is the test error of the model obtained ?
>
> fit.qda <- qda(mpg01 ~ cylinders + weight + displacement + horsepower, data = Auto, subset = train)
> fit.qda
call:
qda(mpg01 ~ cylinders + weight + displacement + horsepower, data = Auto,
    subset = train)

Prior probabilities of groups:
      0      1
0.4571429 0.5428571

Group means:
  cylinders  weight displacement horsepower
0   6.812500 3604.823    271.7396   133.14583
1   4.070175 2314.763    111.6623    77.92105
> |
```

```
> pred.qda <- predict(fit.qda, Auto.test)
> table(pred.qda$class, mpg01.test)
      mpg01.test
      0      1
0  89  13
1  11  69
> |
```

```
> mean(pred.qda$class != mpg01.test)
[1] 0.1318681
> |
```

From above it is clear that test error rate is 13.186%

(f) (5 points) Perform logistic regression on the training data in order to predict mpg01 using the variables that seemed most associated with mpg01 in (b). What is the test error of the model obtained?

```
> #3. f) Perform logistic regression on the training data in order to predict "mpg01" using the variables
> #that seemed most associated with "mpg01" in (b). What is the test error of the model obtained ?
>
> fit.glm <- glm(mpg01 ~ cylinders + weight + displacement + horsepower, data = Auto, family = binomial, subset = train)
> summary(fit.glm)

call:
glm(formula = mpg01 ~ cylinders + weight + displacement + horsepower,
    family = binomial, data = Auto, subset = train)

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-2.48027  -0.03413   0.10583   0.29634   2.57584

Coefficients:
            Estimate Std. Error z value Pr(>|z|)
(Intercept)  17.658730   3.409012   5.180 2.22e-07 ***
cylinders    -1.028032   0.653607  -1.573  0.1158
weight       -0.002922   0.001137  -2.569  0.0102 *
displacement  0.002462   0.015030   0.164  0.8699
horsepower   -0.050611   0.025209  -2.008  0.0447 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

    Null deviance: 289.58  on 209  degrees of freedom
Residual deviance:  83.24  on 205  degrees of freedom
AIC: 93.24

Number of Fisher Scoring iterations: 7
```

```

> probs <- predict(fit.glm, Auto.test, type = "response")
> pred.glm <- rep(0, length(probs))
> pred.glm[probs > 0.5] <- 1
> table(pred.glm, mpg01.test)
      mpg01.test
pred.glm  0  1
      0 89 11
      1 11 71
> |

> mean(pred.glm != mpg01.test)
[1] 0.1208791
> |

```

From above it is clear that test error rate is 12.08791%

(g) (5 points) Perform KNN on the training data, with several values of K, in order to predict mpg01. Use only the variables that seemed most associated with mpg01 in (b). What test errors do you obtain? Which value of K seems to perform the best on this data set?

```

181 #3. g) Perform KNN on the training data, with several values of K, in order to predict "mpg01"
182 #using the variables that seemed most associated with "mpg01" in (b). What test errors do you obtain ?
183 #which value of K seems to perform the best on this data set ?
184 train.X <- cbind(cylinders, weight, displacement, horsepower)[train, ]
185 test.X <- cbind(cylinders, weight, displacement, horsepower)[!train, ]
186 train.mpg01 <- mpg01[train]
187 set.seed(1)
188 pred.knn <- knn(train.X, test.X, train.mpg01, k = 1)
189 table(pred.knn, mpg01.test)
190
179:1 (Top Level) ↕

```

Console	Terminal
<pre> ~ / ↕ mpg01.test pred.knn 0 1 0 83 11 1 17 71 > > mean(pred.knn != mpg01.test) [1] 0.1538462 > </pre>	

For K=1 the test error rate is 15.38462%

```

> pred.knn <- knn(train.X, test.X, train.mpg01, k = 10)
> table(pred.knn, mpg01.test)
      mpg01.test
pred.knn  0  1
      0 77  7
      1 23 75
> |

```

```
> mean(pred.knn != mpg01.test)
[1] 0.1648352
> |
```

For K=10 the test error rate is 16.48352%

```
> pred.knn <- knn(train.X, test.X, train.mpg01, k = 100)
> table(pred.knn, mpg01.test)
      mpg01.test
pred.knn  0   1
         0 81  7
         1 19 75

> mean(pred.knn != mpg01.test)
[1] 0.1428571
> |
```

From above for K=100 the test error rate is 14.28571%