

# Introduction to Statistical Learning Lab4

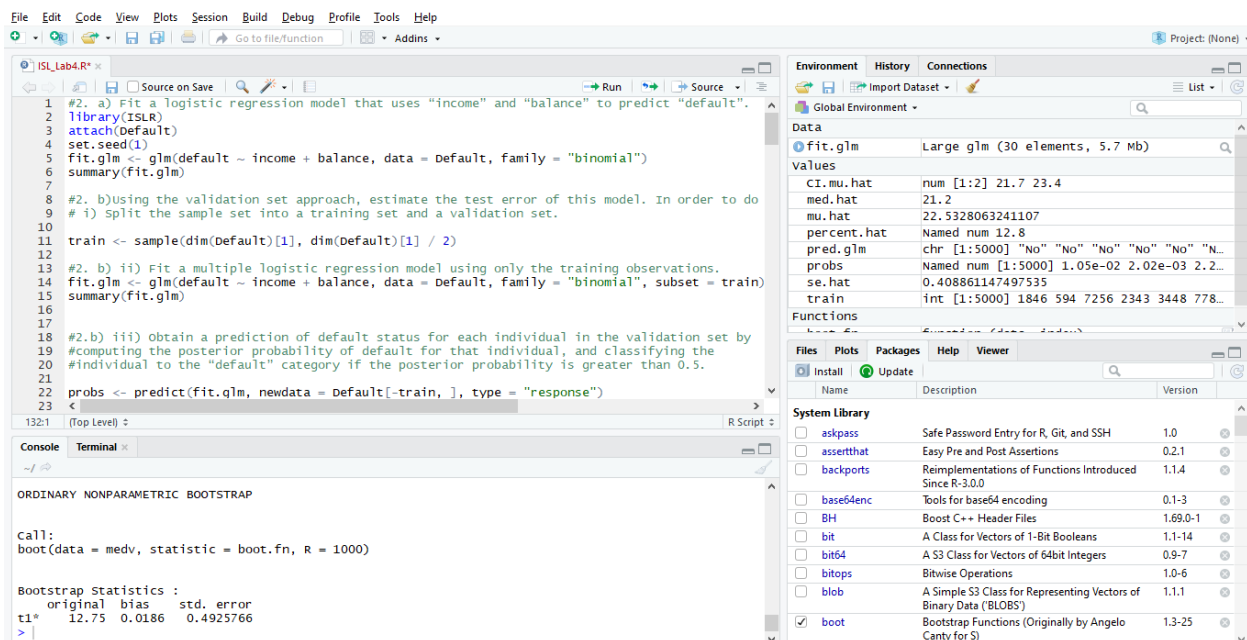
## (Cross Validation and Bootstrap)

Name: Sandeep Reddy Salkuti

Id: 16296868

Email: [sswf7@umsystem.edu](mailto:sswf7@umsystem.edu)

1) You may download the R Code for Labs and the Data Sets to use from the textbook website.



```
1 #2. a) Fit a logistic regression model that uses "income" and "balance" to predict "default".
2 library(ISLR)
3 attach(Default)
4 set.seed(1)
5 fit.glm <- glm(default ~ income + balance, data = Default, family = "binomial")
6 summary(fit.glm)
7
8 #2. b) Using the validation set approach, estimate the test error of this model. In order to do
9 # i) Split the sample set into a training set and a validation set.
10
11 train <- sample(dim(Default)[1], dim(Default)[1] / 2)
12
13 #2. b) ii) Fit a multiple logistic regression model using only the training observations.
14 fit.glm <- glm(default ~ income + balance, data = Default, family = "binomial", subset = train)
15 summary(fit.glm)
16
17 #2. b) iii) Obtain a prediction of default status for each individual in the validation set by
18 # computing the posterior probability of default for that individual, and classifying the
19 # individual to the "default" category if the posterior probability is greater than 0.5.
20
21 probs <- predict(fit.glm, newdata = Default[-train, ], type = "response")
22
23 <
```

Environment

Global Environment

Data

fit.glm Large glm (30 elements, 5.7 Mb)

Values

Variable	Value
ci.mu.hat	num [1:2] 21.7 23.4
med.hat	21.2
mu.hat	22.5328063241107
percent.hat	Named num 12.8
pred.glm	chr [1:5000] "No" "No" "No" "No" "No" "N..."
probs	Named num [1:5000] 1.05e-02 2.02e-03 2.2...
se.hat	0.408861147497535
train	int [1:5000] 1846 594 7256 2343 3448 778...

Functions

Files Plots Packages Help Viewer

System Library

Package	Description	Version
askpass	Safe Password Entry for R, Git, and SSH	1.0
assertthat	Easy Pre and Post Assertions	0.2.1
backports	Reimplementations of Functions Introduced Since R-3.0.0	1.1.4
base64enc	Tools for base64 encoding	0.1-3
BH	Boost C++ Header Files	1.69.0-1
bit	A Class for Vectors of 1-Bit Booleans	1.1-14
bit64	A S3 Class for Vectors of 64bit Integers	0.9-7
bitops	Bitwise Operations	1.0-6
blob	A Simple S3 Class for Representing Vectors of Binary Data ('BLOBs')	1.1.1
boot	Bootstrap Functions (Originally by Angelo Canty for S)	1.3-25

Console

```
ORDINARY NONPARAMETRIC BOOTSTRAP

Call:
boot(data = medv, statistic = boot.fn, R = 1000)

Bootstrap Statistics :
  original bias std. error
t1*    12.75  0.0186   0.4925766
> |
```

2) In Chapter 4, we used logistic regression to predict the probability of default using income and balance on the Default data set. We will now estimate the test error of this logistic regression model using the validation set approach. Do not forget to set a random seed before beginning your analysis.

(a) (10 points) Fit a logistic regression model that uses income and balance to predict default.

```
ISL_Lab4.R x
Source on Save
Run Source
1 #2. a) Fit a logistic regression model that uses "income" and "balance" to predict "default"
2 library(ISLR)
3 attach(Default)
4 set.seed(1)
5 fit.glm <- glm(default ~ income + balance, data = Default, family = "binomial")
6 summary(fit.glm)

1:1 (Top Level) R Script

Console Terminal x
~/
Call:
glm(formula = default ~ income + balance, family = "binomial",
    data = Default)

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-2.4725  -0.1444  -0.0574  -0.0211   3.7245

Coefficients:
            Estimate Std. Error z value Pr(>|z|)
(Intercept) -1.154e+01  4.348e-01 -26.545  < 2e-16 ***
income       2.081e-05  4.985e-06   4.174  2.99e-05 ***
balance      5.647e-03  2.274e-04  24.836  < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

    Null deviance: 2920.6  on 9999  degrees of freedom
Residual deviance: 1579.0  on 9997  degrees of freedom
AIC: 1585

Number of Fisher scoring iterations: 8

> |
```

(b) (10 points total) Using the validation set approach, estimate the test error of this model. In order to do this, you must perform the following steps:

i. (2.5 points) Split the sample set into a training set and a validation set.

```
26 |
27 #2. b) Using the validation set approach, estimate the test error of this model. In order
28 # i) Split the sample set into a training set and a validation set.
29
30 train <- sample(dim(Default)[1], dim(Default)[1] / 2)
31

26:1 (Top Level) R Script

Console Terminal x
~/
> train <- sample(dim(Default)[1], dim(Default)[1] / 2)
. |
```

ii. (2.5 points) Fit a multiple logistic regression model using only the training observations.

```
Console Terminal x
~/
> #2. b) ii) Fit a multiple logistic regression model using only the training observations.
> fit.glm <- glm(default ~ income + balance, data = Default, family = "binomial", subset = train)
> summary(fit.glm)

Call:
glm(formula = default ~ income + balance, family = "binomial",
    data = Default, subset = train)

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-2.5830  -0.1428  -0.0573  -0.0213   3.3395

Coefficients:
            Estimate Std. Error z value Pr(>|z|)
(Intercept) -1.194e+01  6.178e-01 -19.333  < 2e-16 ***
income       3.262e-05  7.024e-06   4.644 3.41e-06 ***
balance      5.689e-03  3.158e-04  18.014  < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

    Null deviance: 1523.8  on 4999  degrees of freedom
Residual deviance:  803.3  on 4997  degrees of freedom
AIC: 809.3

Number of Fisher Scoring iterations: 8

> |
```

iii. (2.5 points) Obtain a prediction of default status for each individual in the validation set by computing the posterior probability of default for that individual, and classifying the individual to the default category if the posterior probability is greater than 0.5.

```
19 |
20 #3.b) iii) Obtain a prediction of default status for each individual in the validation set by
21 #computing the posterior probability of default for that individual, and classifying the
22 #individual to the "default" category if the posterior probability is greater than 0.5.
23
24 probs <- predict(fit.glm, newdata = Default[-train, ], type = "response")
25 pred.glm <- rep("No", length(probs))
26 pred.glm[probs > 0.5] <- "Yes"
27

19:1 (Top Level) R Script
Console Terminal x
> probs <- predict(fit.glm, newdata = Default[-train, ], type = "response")
> pred.glm <- rep("No", length(probs))
> pred.glm[probs > 0.5] <- "Yes"
> |
```



```

43
44 train <- sample(dim(Default)[1], dim(Default)[1] / 2)
45 fit.glm <- glm(default ~ income + balance, data = Default, family = "binomial", subset = train)
46 probs <- predict(fit.glm, newdata = Default[-train, ], type = "response")
47 pred.glm <- rep("No", length(probs))
48 pred.glm[probs > 0.5] <- "Yes"
49 mean(pred.glm != Default[-train, ]$default)
50 <
42:1 (Top Level) ↕

```

Console Terminal ×

```

~/
> train <- sample(dim(Default)[1], dim(Default)[1] / 2)
> fit.glm <- glm(default ~ income + balance, data = Default, family = "binomial", subset = train)
> probs <- predict(fit.glm, newdata = Default[-train, ], type = "response")
> pred.glm <- rep("No", length(probs))
> pred.glm[probs > 0.5] <- "Yes"
> mean(pred.glm != Default[-train, ]$default)
[1] 0.0244
>

```

```

46
47 train <- sample(dim(Default)[1], dim(Default)[1] / 2)
48 fit.glm <- glm(default ~ income + balance, data = Default, family = "binomial", subset = train)
49 probs <- predict(fit.glm, newdata = Default[-train, ], type = "response")
50 pred.glm <- rep("No", length(probs))
51 pred.glm[probs > 0.5] <- "Yes"
52 mean(pred.glm != Default[-train, ]$default)
53 <
47:1 (Top Level) ↕

```

Console Terminal ×

```

~/
> train <- sample(dim(Default)[1], dim(Default)[1] / 2)
> fit.glm <- glm(default ~ income + balance, data = Default, family = "binomial", subset = train)
> probs <- predict(fit.glm, newdata = Default[-train, ], type = "response")
> pred.glm <- rep("No", length(probs))
> pred.glm[probs > 0.5] <- "Yes"
> mean(pred.glm != Default[-train, ]$default)
[1] 0.0244
>

```

From above it is clear that validation test error rate is varying depending on which observations are included in training set and which observations are included in validation set.

**(d) (10 points)** Now consider a logistic regression model that predicts the probability of default using income, balance, and a dummy variable for student. Estimate the test error for this model using the validation set approach. Comment on whether or not including a dummy variable for student leads to a reduction in the test error rate.

```

54 #2. d) Now consider a logistic regression model that predicts the probability of "default" using
55 # "income", "balance", and a dummy variable for "student". Estimate the test error for this model
56 # using the validation set approach. Comment on whether or not including a dummy variable for "student"
57 # leads to a reduction in the test error rate.
58
59 train <- sample(dim(Default)[1], dim(Default)[1] / 2)
60 fit.glm <- glm(default ~ income + balance + student, data = Default, family = "binomial", subset = train)
61 pred.glm <- rep("No", length(probs))
62 probs <- predict(fit.glm, newdata = Default[-train, ], type = "response")
63 pred.glm[probs > 0.5] <- "Yes"
64 mean(pred.glm != Default[-train, ]$default)
65 <
66
54:1 (Top Level) ↕ R 5

```

---

```

Console Terminal x
~/
> train <- sample(dim(Default)[1], dim(Default)[1] / 2)
> fit.glm <- glm(default ~ income + balance + student, data = Default, family = "binomial", subset = train)
> pred.glm <- rep("No", length(probs))
> probs <- predict(fit.glm, newdata = Default[-train, ], type = "response")
> pred.glm[probs > 0.5] <- "Yes"
> mean(pred.glm != Default[-train, ]$default)
[1] 0.0278
>

```

From above it doesn't seem that by adding "student" dummy variable leads to reduction in validation set estimate of the test error rate.

**3. (40 points)** We continue to consider the use of a logistic regression model to predict the probability of default using income and balance on the Default data set. In particular, we will now compute estimates for the standard errors of the income and balance logistic regression coefficients in two different ways: (1) using the bootstrap, and (2) using the standard formula for computing the standard errors in the `glm()` function. Do not forget to set a random seed before beginning your analysis.

**(a) (10 points)** Using the `summary()` and `glm()` functions, determine the estimated standard errors for the coefficients associated with income and balance in a multiple logistic regression model that uses both predictors.

```

> set.seed(1)
> attach(Default)
The following objects are masked from Default (pos = 3):

    balance, default, income, student

The following objects are masked from Default (pos = 4):

    balance, default, income, student

> fit.glm <- glm(default ~ income + balance, data = Default, family = "binomial")
> summary(fit.glm)

Call:
glm(formula = default ~ income + balance, family = "binomial",
    data = Default)

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-2.4725  -0.1444  -0.0574  -0.0211   3.7245

Coefficients:
            Estimate Std. Error z value Pr(>|z|)
(Intercept) -1.154e+01  4.348e-01 -26.545  < 2e-16 ***
income       2.081e-05  4.985e-06   4.174 2.99e-05 ***
balance      5.647e-03  2.274e-04  24.836  < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

    Null deviance: 2920.6  on 9999  degrees of freedom
Residual deviance: 1579.0  on 9997  degrees of freedom
AIC: 1585

Number of Fisher Scoring iterations: 8

```

The glm() estimates of the standard errors for the coefficients  $\beta_0$ ,  $\beta_1$  and  $\beta_2$  are respectively 0.4347564,  $4.9851672 \times 10^{-6}$  and  $2.2737314 \times 10^{-4}$ .

**(b) (10 points) Write a function, boot.fn(), that takes as input the Default data set as well as an index of the observations, and that outputs the coefficient estimates for income and balance in the multiple logistic regression model.**

```

89 |
90 | #3. b) write a function, boot.fn(), that takes as input the "Default" data set as well as an index of the
91 | #observations, and that outputs the coefficient estimates for "income" and "balance" in the multiple logistic regression model.
92 |
93 |
94 | boot.fn <- function(data, index) {
95 |   fit <- glm(default ~ income + balance, data = data, family = "binomial", subset = index)
96 |   return (coef(fit))
97 | }
98 |
99 |
100 |

```

90:1 (Top Level) ↕ R Script

Console Terminal

```

> boot.fn <- function(data, index) {
+   fit <- glm(default ~ income + balance, data = data, family = "binomial", subset = index)
+   return (coef(fit))
+ }
>

```

- (c) (10 points) Use the `boot()` function together with your `boot.fn()` function to estimate the standard errors of the logistic regression coefficients for income and balance.

```
82 #3. c) Use the boot() function together with your boot.fn() function to estimate the standard errors of the logistic
83 #regression coefficients for "income" and "balance".
84
85 library(boot)
86 boot(Default, boot.fn, 1000)
87
88
89
90
```

82:1 (Top Level) R Script

Console Terminal

```
~/
package 'boot' successfully unpacked and MD5 sums checked

The downloaded binary packages are in
  C:\Users\SandeepReddy\AppData\Local\Temp\Rtmp8CuYvp\downloaded_packages
> library(boot)
warning message:
package 'boot' was built under R version 3.6.3
> boot(Default, boot.fn, 1000)

ORDINARY NONPARAMETRIC BOOTSTRAP

Call:
boot(data = Default, statistic = boot.fn, R = 1000)

Bootstrap Statistics :
      original      bias      std. error
t1*  -1.154047e+01  -3.945460e-02  4.344722e-01
t2*   2.080898e-05  1.680317e-07  4.866284e-06
t3*   5.647103e-03  1.855765e-05  2.298949e-04
```

The bootstrap estimates of the standard errors for the coefficients  $\beta_0$ ,  $\beta_1$  and  $\beta_2$  are respectively 0.4344,  $4.8662 \times 10^{-6}$  and  $2.2989 \times 10^{-4}$ .

- (d) (10 points) Comment on the estimated standard errors obtained using the `glm()` function and using your bootstrap function.

From the above observations we can assume that, the estimated standard errors obtained by the two methods are close.

#### 4. (40 points) We will now consider the Boston housing data set, from the MASS library.

- (a) (5 points) Based on this data set, provide an estimate for the population mean of `medv`. Call this estimate  $\hat{\mu}$ .

```
106 |
107 #4. a) Based on this data set, provide an estimate for the population mean of "medv". Call this estimate  $\hat{\mu}$ .
108 library(MASS)
109 attach(Boston)
110 mu.hat <- mean(medv)
111 mu.hat
112 <
```

106:1 (Top Level)

Console Terminal

```
~/
[1] 22.53281
>
```



(b) Provide an estimate of the standard error of  $\hat{\mu}$ . Interpret this result. Hint: We can compute the standard error of the sample mean by dividing the sample standard deviation by the square root of the number of observations.

```

107 |
108 |
109 #4. b)Provide an estimate of the standard error of  $\mu^{\wedge}$ . Interpret this result.
110 se.hat <- sd(medv) / sqrt(dim(Boston)[1])
111 se.hat
112 <

```

108:1 (Top Level) ↕

Console Terminal ×

~/

```

[1] 0.4088611
> |

```

(c) (5 points) Now estimate the standard error of  $\hat{\mu}$  using the bootstrap. How does this compare to your answer from (b)?

```

97 |
98 #4. c)Now estimate the standard error of  $\mu^{\wedge}$  using the bootstrap. How does this compare to your answer from (b) ?
99 set.seed(1)
100 boot.fn <- function(data, index) {
101   mu <- mean(data[index])
102   return (mu)
103 }
104 boot(medv, boot.fn, 1000)
105
106 <
107

```

97:1 (Top Level) ↕ R Sc

Console Terminal ×

~/

```

ORDINARY NONPARAMETRIC BOOTSTRAP

Call:
boot(data = medv, statistic = boot.fn, R = 1000)

Bootstrap Statistics :
  original    bias   std. error
t1*  22.53281 0.007650791   0.4106622
> |

```

The bootstrap estimated standard error of  $\mu^{\wedge}$  of 0.4106 is very close to the estimate found in (b) of 0.4089.

(d) (5 points) Based on your bootstrap estimate from (c), provide a 95% confidence interval for the mean of medv. Compare it to the results obtained using `t.test(Boston$medv)`.

```

111 #4. d)Based on your bootstrap estimate from (c), provide a 95% confidence interval for the mean of "medv".
112 #Compare it to the results obtained using t.test(Boston$medv).
113
114 t.test(medv)
115
116 CI.mu.hat <- c(22.53 - 2 * 0.4119, 22.53 + 2 * 0.4119)
117 CI.mu.hat
118
119 <
120
121 111:1 (Top Level) ⚡

```

---

```

Console Terminal ✕
~/
> t.test(medv)

One Sample t-test

data: medv
t = 55.111, df = 505, p-value < 2.2e-16
alternative hypothesis: true mean is not equal to 0
95 percent confidence interval:
 21.72953 23.33608
sample estimates:
mean of x
 22.53281

> CI.mu.hat <- c(22.53 - 2 * 0.4119, 22.53 + 2 * 0.4119)
> CI.mu.hat
[1] 21.7062 23.3538
> |

```

The bootstrap confidence interval is very close to the one provided by the `t.test()` function.

**(e) (5 points) Based on this data set, provide an estimate,  $\hat{\mu}_{med}$ , for the median value of medv in the population.**

```

127
128 #4. e)Based on this data set, provide an estimate,  $\mu_{med}$ , for the median value of "medv" in the population.
129 med.hat <- median(medv)
130 med.hat
131
132 <
133
134 127:1 (Top Level) ⚡

```

---

```

Console Terminal ✕
~/
> med.hat <- median(medv)
> med.hat
[1] 21.2
> |

```

- (f) (5 points) We now would like to estimate the standard error of  $\hat{\mu}_{med}$ . Unfortunately, there is no simple formula for computing the standard error of the median. Instead, estimate the standard error of the median using the bootstrap. Comment on your findings.

```
118 #4. f) We now would like to estimate the standard error of  $\hat{\mu}_{med}$ . Unfortunately, there is no simple formula for
119 #computing the standard error of the median. Instead, estimate the standard error of the median using the
120 #bootstrap. Comment on your findings.
121
122 boot.fn <- function(data, index) {
123   mu <- median(data[index])
124   return(mu)
125 }
126 boot(medv, boot.fn, 1000)
127 <
128
118:1 (Top Level) ↕
```

Console Terminal x

```
~/ |
> boot.fn <- function(data, index) {
+   mu <- median(data[index])
+   return(mu)
+ }
> boot(medv, boot.fn, 1000)

ORDINARY NONPARAMETRIC BOOTSTRAP

Call:
boot(data = medv, statistic = boot.fn, R = 1000)

Bootstrap Statistics :
      original    bias      std. error
t1*      21.2 -0.0386    0.3770241
> |
```

From above we get estimated median value of 21.2 which is equal to the value obtained in (e), with a standard error of 0.37702 which is relatively small compared to median value.

- (g) Based on this data set, provide an estimate for the tenth percentile of medv in Boston suburbs. Call this quantity  $\hat{\mu}_{0.1}$  (You can use the quantile() function.)

```
139 |
140 #4. g) Based on this data set, provide an estimate for the tenth percentile of "medv" in Boston suburbs.
141 #Call this quantity  $\hat{\mu}_{0.1}$ .
142 percent.hat <- quantile(medv, c(0.1))
143 percent.hat
144 <
139:1 (Top Level) ↕
```

Console Terminal x

```
~/ |
> percent.hat <- quantile(medv, c(0.1))
> percent.hat
10%
12.75
> |
```

(h) (5 points) Use the bootstrap to estimate the standard error of  $\hat{\mu}_{med}$ . Comment on your findings.

```
140 #4. h)Use the bootstrap to estimate the standard error of  $\mu^{0.1}$ . Comment on your findings.
141
142 boot.fn <- function(data, index) {
143   mu <- quantile(data[index], c(0.1))
144   return (mu)
145 }
146 boot(medv, boot.fn, 1000)
147 <
```

138:1 (Top Level) ▾

Console Terminal ×

~/ | ↻

ORDINARY NONPARAMETRIC BOOTSTRAP

Call:  
boot(data = medv, statistic = boot.fn, R = 1000)

Bootstrap Statistics :

	original	bias	std. error
t1*	12.75	0.0186	0.4925766

> |

From above we get an estimated 10% value of 12.75 which is equal to the value obtained in (g), with a standard error of 0.4925 which is relatively small compared to percentile value.